
Tools, Not Talk: The Knowledge Ceiling of Biology Agents

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Agentic systems for biology are increasingly built around orchestration: critics
2 that review answers, teams of role-playing scientists that debate, and models that
3 reason for longer. We show that on objectively scored biology research tasks this
4 machinery adds little beyond what matched independent sampling provides, and
5 we identify why. Across 150 pre-registered LAB-Bench and LABBench2 ques-
6 tions, compared at a matched budget of three model calls, self-critique never beats
7 simple majority voting, multi-agent debate helps only a model whose reasoning
8 is switched off, and a 25-fold increase in generated tokens from longer reasoning
9 does not improve accuracy. The reason is a ceiling set by the questions them-
10 selves: which question is asked explains 56% of the variance in correctness, while
11 the choice of model, strategy and reasoning effort explains 0.6%. That ceiling is
12 largely missing knowledge rather than missing deliberation. A single model call
13 with web search and code execution solves 47% of the questions that all eight
14 orchestrated systems in our main study failed. For life-science agents, the budget
15 belongs in tools, not talk.

16 1 Introduction

17 LLM agents now plan experiments, analyse data and generate hypotheses in chemistry and biology
18 [1, 2, 5, 6, 8, 16]. Many of these systems are organised as conversations: a critic reviews a draft,
19 several specialised scientists debate, an adjudicator decides. Their gains are usually reported against
20 a single model call, even though the agent makes many. Sampling several answers and voting already
21 improves accuracy [3, 12, 19], so these comparisons cannot tell whether the orchestration helps or
22 whether the extra calls do [9].

23 We separate the two. Every orchestrated strategy receives exactly the budget of a voting baseline:
24 three calls, the same model, the same prompts and the same decoding. We then ask what actually
25 moves accuracy on biology research tasks: orchestration, longer reasoning, or access to knowledge
26 through tools.

27 Contributions.

- 28 • **Orchestration adds little at matched compute.** Across three models, self-critique never
29 beats voting and repairs no more answers than it breaks. Debate helps only when the
30 model cannot reason on its own. Raising reasoning effort, for six model families, rarely
31 helps either (§3.1–3.2).
- 32 • **The question decides.** Every system solves the same questions and fails the same ques-
33 tions. Question identity explains 56% of the variance in correctness and the system 0.6%
34 (§3.3).

Table 1: Accuracy (%) at a matched budget of three calls (DIRECT uses one). Δ = DEBATE – VOTE in points, with paired bootstrap 95% CI.

Model	n	DIRECT	VOTE	CRITIQUE	DEBATE	Δ [95% CI]
GPT-OSS-120B	150	46.7	49.3	42.7	47.3	-2.0 [-8.7, +4.7]
Gemma 4 31B (no thinking)	150	48.7	48.7	47.3	57.3	+8.7 [+2.0, +15.3]
Flash-Lite	75	61.3	62.7	56.0	57.3	-5.3 [-13.3, +1.3]

- **Tools break the ceiling.** One tool-using call solves 47% of the questions that no main-study system solved, and more compute on top of tools adds comparatively little (§3.4).

2 Setup

Tasks. We draw 30 questions from each of five categories (150 in total): four from LAB-Bench [10] (protocol troubleshooting, sequence manipulation, database and literature questions; multiple choice) and open-ended sequence computations from LABBench2 SeqQA2 [11], scored with the official validators. All answers are scored deterministically, with no LLM judge.

Strategies. DIRECT makes one call. VOTE makes three independent calls and takes the majority answer. CRITIQUE solves, critiques its own solution, then revises. DEBATE has a molecular biologist and a skeptical reviewer solve independently, then an adjudicator decides. VOTE, CRITIQUE and DEBATE each use exactly three calls, and only visible outputs pass between calls (prompts in App. A).

Models. The matched-budget study uses GPT-OSS-120B [13] (low reasoning effort), Gemma 4 31B [17] (thinking disabled) and Gemini 3.5 Flash-Lite (low thinking) on all 150 questions (Flash-Lite on a fixed 75). To test reasoning effort and tools, we add GPT-6 Luna, GPT-6 Sol and GPT-5.6 Terra, and rerun GPT-OSS and Flash-Lite at high effort, on a fixed 30-question subset with 6 per category. Task IDs, prompts, decoding and the primary analysis were frozen before evaluation (App. B). All contrasts are paired over questions, with bootstrap 95% CIs and exact McNemar tests.

3 Results

3.1 Orchestration adds little at matched compute

Three calls spent on orchestration generally do no better than three calls spent on independent samples (Table 1). **Self-critique does not beat voting** for any of the three models (−6.7, −1.3, −6.7 points; pooled −4.5 [−9.3, +0.3]). It breaks correct answers as readily as it fixes wrong ones: relative to the vote it rescues 24%, 13% and 25% of errors but corrupts 38%, 16% and 26% of correct answers. The damage is worst on exact sequence manipulation (SeqQA: −23 and −20 points for GPT-OSS and Flash-Lite), where a fluent but wrong critique argues the model out of a correct computation. **Debate helps only the model that cannot think.** For Gemma, whose thinking is disabled, debate beats voting by +8.7 points ($p = 0.015$); for the two models that reason, it does not (−2.0 and −5.3).

3.2 Longer reasoning does not substitute for knowledge

Raising reasoning effort multiplies the cost without moving accuracy (Fig. 1). GPT-OSS at high effort generates $25\times$ more tokens and scores 47% against 50%; Flash-Lite spends $5\times$ more tokens and scores 67% against 70%. The one model that gains is GPT-6 Luna, which uses no reasoning tokens at low effort, and it gains only up to the level its low-effort debate had already reached (60% vs. 60%). Debate and reasoning are therefore substitutes, and neither gets past the plateau. Stacking orchestration on high effort adds little: GPT-6 Sol scores 70% with every strategy while the three-call strategies generate $\sim 3\times$ the tokens.

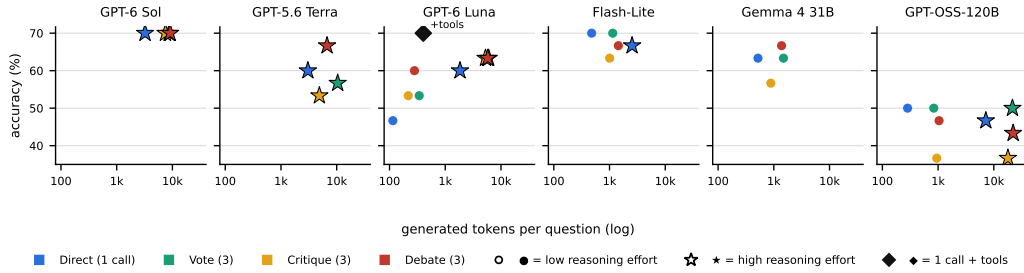


Figure 1: Every way of spending inference lands on a similar per-model plateau; tools rise above it. Accuracy vs. generated tokens per question on a shared 30-question set, one panel per model. Colour marks the strategy. Marker shape marks reasoning effort: circles (●) are low effort and stars (★) high effort. GPT-6 Sol and GPT-5.6 Terra were run at high effort only. Across nearly two orders of magnitude in tokens, accuracy moves little within a model. The black diamond (◆) is one GPT-6 Luna call with web search and code execution. It is the only configuration clearly above its model’s plateau, and among the cheapest.

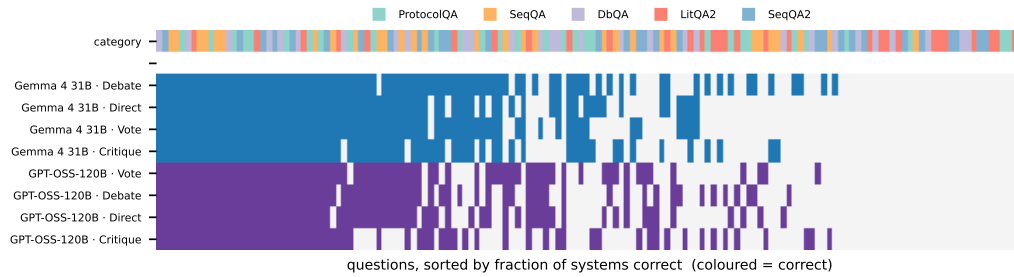


Figure 2: The question, far more than the system, decides. Correctness of every model $\times$ strategy system (rows) on every question (columns, sorted by the fraction of systems correct). Systems largely solve and fail the same questions. Question identity explains 56% [49, 62] of the variance in correctness; the system explains 0.6% [0.3, 1.9].

73 3.3 The question, not the system, decides

74 A two-way decomposition of the correctness matrix attributes 56% of the variance to the question
 75 and 0.6% to the system (Fig. 2). The pattern holds across all 33 model, strategy and effort con-
 76 figurations on the shared subset, frontier models at high effort included (47% vs. 3.9%). Systems
 77 largely solve the same easy questions and fail the same hard ones. The 32 questions that no system
 78 solves are dominated by literature and database recall (18) and open-ended sequence computation
 79 (9). These are failures of knowledge, which additional deliberation did not supply in any configura-
 80 tion we tested.

81 3.4 Tools break the ceiling

82 Giving GPT-6 Luna web search and code execution in a single call raises accuracy on the shared
 83 subset from 47% to 70% (+23 points [+7, +43], $p = 0.039$). That is more than any orchestration
 84 or reasoning setting achieves, at ~ 404 generated tokens per question. On the 32 questions that all
 85 eight main-study systems (two models $\times$ four strategies) failed, the same single call reaches 47%,
 86 against 6% without tools (+41 points [+25, +56], $p = 0.0002$). Literature questions rise from 18%
 87 to 91% and database questions from 0% to 43%. Retrieval helping on these categories is expected,
 88 since they were designed for retrieval agents [10, 14]. The point is the contrast: none of the eight
 89 orchestrated systems solved any of them. Compute on top of tools adds comparatively little. High
 90 effort reaches 83% (+13, $p = 0.22$) and a three-call vote 77% (+7, $p = 0.73$), at $\sim 6\times$ and $\sim 3.5\times$
 91 the tokens. Tool calls never reached the benchmark or its datasets.

92 4 Implications for building life-science agents

93 **Benchmark against a matched vote.** A multi-call agent should beat independent sampling with the
94 same number of calls, not a single call [9]. Most orchestration gains disappear under this baseline,
95 consistent with compute-matched studies of debate in general reasoning [4, 15, 18, 20] and with the
96 unreliability of intrinsic self-correction [7]. **Spend the budget on tools, not talk.** Critics, debaters
97 and longer chains of thought re-deliberate over knowledge the model does not have. Retrieval sup-
98 plies it, at a fraction of the tokens. **Build domain tools for computation.** A generic shell barely
99 helps open-ended sequence design and analysis (0% to 11%). The next gains lie in biology-specific
100 tool ecosystems and specialist models, not in more conversation between copies of the same model.

101 **Scope.** Our tasks have deterministic answers, and we study short-horizon question answering
102 rather than long-running discovery loops. The frontier and tool comparisons use a 30-question
103 subset, and the tool condition adds a one-sentence note announcing the tools.

104 5 Conclusion

105 On biology research tasks, how a model spends extra inference barely changes which questions it
106 gets right. This holds for voting, critique, debate and longer reasoning alike: the question explains
107 56% of the variance in correctness and the system 0.6%. The ceiling is largely knowledge, and
108 a single tool-using call lifts it where orchestration did not. Life-science agents should spend their
109 budget on tools, not talk.

110 References

- 111 [1] Daniil A. Boiko, Robert MacKnight, Ben Kline, and Gabe Gomes. Autonomous chemical
112 research with large language models. *Nature*, 624:570–578, 2023. URL [https://doi.org/](https://doi.org/10.1038/s41586-023-06792-0)
113 [10.1038/s41586-023-06792-0](https://doi.org/10.1038/s41586-023-06792-0).
- 114 [2] Andres M. Bran, Sam Cox, Andrew D. White, and Philippe Schwaller. Chemcrow: Augment-
115 ing large-language models with chemistry tools. *arXiv preprint arXiv:2304.05376*, 2023. URL
116 <https://arxiv.org/abs/2304.05376>.
- 117 [3] Bradley Brown, Jordan Juravsky, Ryan Ehrlich, Ronald Clark, Quoc V. Le, Christopher Ré,
118 and Azalia Mirhoseini. Large language monkeys: Scaling inference compute with repeated
119 sampling. *arXiv preprint arXiv:2407.21787*, 2024. URL [https://arxiv.org/abs/2407.](https://arxiv.org/abs/2407.21787)
120 [21787](https://arxiv.org/abs/2407.21787).
- 121 [4] H. K. Choi, X. Zhu, and S. Li. Debate or vote: Which yields better decisions in multi-agent
122 large language models? *arXiv preprint arXiv:2508.17536*, 2025. URL [https://arxiv.org/](https://arxiv.org/abs/2508.17536)
123 [abs/2508.17536](https://arxiv.org/abs/2508.17536). NeurIPS 2025 Spotlight.
- 124 [5] Ali E. Ghareeb, Benjamin Chang, Ludovico Mitchener, Angela Yiu, Caralyn J. Szostkiewicz,
125 Dmytro Shved, Gavin J. Gyimesi, Jon M. Laurent, et al. A multi-agent system for
126 automating scientific discovery. *Nature*, 2026. URL [https://doi.org/10.1038/](https://doi.org/10.1038/s41586-026-10652-y)
127 [s41586-026-10652-y](https://doi.org/10.1038/s41586-026-10652-y).
- 128 [6] Juraj Gottweis, Wei-Hung Weng, Alexander Daryin, Tao Tu, Petar Sirkovic, Artiom
129 Myaskovsky, Grzegorz Glowaty, Felix Weissenberger, et al. Accelerating scientific discovery
130 with co-scientist. *Nature*, 2026. URL <https://doi.org/10.1038/s41586-026-10644-y>.
- 131 [7] Jie Huang, Xinyun Chen, Swaroop Mishra, Huaixiu Steven Zheng, Adams Wei Yu, Xinying
132 Song, and Denny Zhou. Large language models cannot self-correct reasoning yet. *arXiv*
133 *preprint arXiv:2310.01798*, 2023. URL <https://arxiv.org/abs/2310.01798>. ICLR
134 2024.
- 135 [8] Kexin Huang, Serena Zhang, Hanchen Wang, Yuanhao Qu, Yingzhou Lu, Yusuf Roohani,
136 Ryan Li, Lin Qiu, et al. Biomni: A general-purpose biomedical ai agent. *bioRxiv preprint*,
137 2025. URL <https://www.biorxiv.org/content/10.1101/2025.05.30.656746v1>.
- 138 [9] Sayash Kapoor, Benedikt Stroebel, Zachary S. Siegel, Nitya Nadgir, and Arvind Narayanan. Ai
139 agents that matter. *arXiv preprint arXiv:2407.01502*, 2024. URL [https://arxiv.org/abs/](https://arxiv.org/abs/2407.01502)
140 [2407.01502](https://arxiv.org/abs/2407.01502).

- [10] Jon M. Laurent, Joseph D. Janizek, Michael Ruzo, Michaela M. Hinks, Michael J. Hammerling, Siddharth Narayanan, Manvitha Ponnampati, Andrew D. White, and Samuel G. Rodrigues. Lab-bench: Measuring capabilities of language models for biology research. *arXiv preprint arXiv:2407.10362*, 2024. URL <https://arxiv.org/abs/2407.10362>.
- [11] Jon M Laurent, Albert Bou, Michael Pieler, Conor Igoe, Alex Andonian, Siddharth Narayanan, James Braza, Alexandros Sanchez Vassopoulos, et al. Labbench2: An improved benchmark for ai systems performing biology research. *arXiv preprint arXiv:2604.09554*, 2026. URL <https://arxiv.org/abs/2604.09554>.
- [12] Lingyao Li and Ze Zhang. More agents is all you need. *Transactions on Machine Learning Research*, 2024. URL <https://arxiv.org/abs/2402.05120>.
- [13] OpenAI. gpt-oss-120b & gpt-oss-20b model card. *arXiv preprint arXiv:2508.10925*, 2025. URL <https://arxiv.org/abs/2508.10925>.
- [14] Michael D. Skarlinski, Sam Cox, Jon M. Laurent, James D. Braza, Michaela Hinks, Michael J. Hammerling, Manvitha Ponnampati, Samuel G. Rodrigues, et al. Language agents achieve superhuman synthesis of scientific knowledge. *arXiv preprint arXiv:2409.13740*, 2024. URL <https://arxiv.org/abs/2409.13740>.
- [15] Andries Smit, Nathan Grinsztajn, Paul Duckworth, Thomas D. Barrett, and Arnau Pretorius. Should we be going mad? a look at multi-agent debate strategies for llms. In *International Conference on Machine Learning*, 2024. URL <https://arxiv.org/abs/2311.17371>.
- [16] Kyle Swanson, Wesley Wu, Nash L. Bulaong, John E. Pak, and James Zou. The virtual lab: Ai agents design new sars-cov-2 nanobodies with experimental validation. *bioRxiv preprint*, 2024. URL <https://www.biorxiv.org/content/10.1101/2024.11.11.623004v1>. Published in Nature 2025.
- [17] Gemma Team, Sherif El Abd, Vaibhav Aggarwal, Robin Algayres, Alek Andreev, Olivier Bachem, Ian Ballantyne, Cormac Brick, et al. Gemma 4 technical report. *arXiv preprint arXiv:2607.02770*, 2026. URL <https://arxiv.org/abs/2607.02770>.
- [18] Qineng Wang, Zihao Wang, Ying Su, Hanghang Tong, and Yangqiu Song. Rethinking the bounds of llm reasoning: Are multi-agent discussions the key? *arXiv preprint arXiv:2402.18272*, 2024. URL <https://arxiv.org/abs/2402.18272>.
- [19] Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-consistency improves chain of thought reasoning in language models. *arXiv preprint arXiv:2203.11171*, 2022. URL <https://arxiv.org/abs/2203.11171>.
- [20] Haolun Wu, Zhenkun Li, and Lingyao Li. Can llm agents really debate? a controlled study of multi-agent debate in logical reasoning. *arXiv preprint arXiv:2511.07784*, 2025. URL <https://arxiv.org/abs/2511.07784>.

A Prompts

All prompts are single user messages. {QUESTION} is the question (with protocol, options or injected files) and {ANSWER INSTRUCTION} is the format line for multiple choice or the official SeqQA2 suffix.

Direct / Independent×3 / Critique call 1 (solve)

Solve the following biology research question.

{QUESTION}

{ANSWER INSTRUCTION}

Respond in exactly this format and keep it short:

<answer>YOUR ANSWER</answer>

CONFIDENCE: <probability between 0 and 1 that your answer is correct>

JUSTIFICATION: <at most two sentences>

192 **Critique call 2 (critique)**

193 Below are a biology research question and a proposed solution. Act as a critical reviewer: identify any
194 scientific, computational or reasoning error in the solution. If you think the answer is wrong, say which
195 answer you believe is correct and why. If you find no error, say so.

196
197 {QUESTION}

198
199 Proposed solution:
200 {CALL 1 OUTPUT}

201
202 Respond with at most four sentences, starting with "CRITIQUE:".

203 **Critique call 3 (revise)**

204 Below are a biology research question, a first attempt at solving it, and a reviewer's critique of that
205 attempt. Either may contain mistakes. Considering both, produce the final answer.

206
207 {QUESTION}

208
209 First attempt:
210 {CALL 1 OUTPUT}

211
212 Reviewer's critique:
213 {CALL 2 OUTPUT}

214
215 {ANSWER INSTRUCTION}

216
217 Respond in exactly this format and keep it short:
218 <answer>YOUR ANSWER</answer>
219 CONFIDENCE: <probability between 0 and 1 that your answer is correct>
220 JUSTIFICATION: <at most two sentences>

221 **Debate call 1 (researcher)**

222 You are a molecular biology researcher. Solve the following biology research question independently.

223
224 {QUESTION}

225
226 {ANSWER INSTRUCTION}

227
228 Respond in exactly this format and keep it short:
229 <answer>YOUR ANSWER</answer>
230 CONFIDENCE: <probability between 0 and 1 that your answer is correct>
231 JUSTIFICATION: <at most two sentences>

232 **Debate call 2 (skeptic)**

233 You are a skeptical scientific reviewer. Solve the following biology research question independently, looking
234 especially for subtle experimental, biological or computational errors and traps.

235
236 {QUESTION}

237
238 {ANSWER INSTRUCTION}

239
240 Respond in exactly this format and keep it short:
241 <answer>YOUR ANSWER</answer>
242 CONFIDENCE: <probability between 0 and 1 that your answer is correct>
243 JUSTIFICATION: <at most two sentences>

244 **Debate call 3 (adjudicator)**

245 Two scientists independently analysed the biology research question below. Adjudicate between their analyses
246 (either or both may be wrong) and submit the final answer.

247
248 {QUESTION}

249

250 Analysis 1 (molecular biology researcher):
 251 {CALL 1 OUTPUT}
 252
 253 Analysis 2 (skeptical scientific reviewer):
 254 {CALL 2 OUTPUT}
 255
 256 {ANSWER INSTRUCTION}
 257
 258 Respond in exactly this format and keep it short:
 259 <answer>YOUR ANSWER</answer>
 260 CONFIDENCE: <probability between 0 and 1 that your answer is correct>
 261 JUSTIFICATION: <at most two sentences>

262 B Pre-registration deviations

263 All task IDs, prompts, decoding settings and the primary analysis were frozen after a 10-question
 264 pilot and before any main-set call. Changes made after the freeze: (i) throttle and retry fixes for an
 265 undocumented 150 requests/hour provider limit (no completed call was re-run); (ii) 22 GPT-OSS
 266 questions served in full by a second provider of the same model (identical prompts, seeds and cap;
 267 excluding them leaves DEBATE – VOTE at -1.6), with partially served questions discarded and re-
 268 run on the primary provider; (iii) the reasoning-effort and tool arms on the fixed 30-question subset
 269 and the unsolved set, added after the freeze. Frontier models were served through their vendor CLI
 270 with web search, code execution and sub-agents disabled except in the tool condition (a first attempt
 271 with web search left on was discarded unanalysed); GPT-OSS at high effort was served by a second
 272 host whose low-effort accuracy matches the main run. The tool condition prepends one sentence
 273 announcing the tools, and all searches and commands were logged and checked for benchmark
 274 access (none).

275 C Additional results

Category	Gemma					Lite					OSS				
	D	I	C	M	Disagree	D	I	C	M	Disagree	D	I	C	M	Disagree
DbQA	56.7	53.3	40.0	56.7	33.3	40.0	40.0	40.0	33.3	26.7	46.7	46.7	43.3	36.7	36.7
LitQA2	33.3	30.0	33.3	36.7	46.7	60.0	53.3	53.3	46.7	26.7	20.0	30.0	40.0	36.7	43.3
ProtocolQA	56.7	56.7	60.0	70.0	20.0	66.7	80.0	86.7	73.3	26.7	56.7	50.0	43.3	46.7	33.3
SeqQA	56.7	56.7	66.7	73.3	46.7	73.3	73.3	53.3	66.7	13.3	63.3	66.7	43.3	60.0	33.3
SeqQA2	40.0	46.7	36.7	50.0	63.3	66.7	66.7	46.7	66.7	40.0	46.7	53.3	43.3	56.7	43.3

Table 2: Per-category accuracy (D=Direct, I=Independent×3, C=Critique, M=Multi). Numbers are percentages.