

# Inference Can’t Buy Knowledge: Agent Orchestration and Extended Reasoning Hit the Same Ceiling on Biology Research Tasks

Anonymous Authors

## Abstract

Agentic systems for biology usually report gains over a single model call while spending several calls. We ask where such gains come from on objectively scored biology research tasks (LAB-Bench and LABBench2: protocol troubleshooting, sequence manipulation, database and literature questions). At a matched budget of three calls, we compare independent sampling with majority vote, sequential self-critique, and a two-scientist debate with an adjudicator, across three open and API models on 150 pre-registered questions scored without an LLM judge, and we contrast these with simply letting frontier models reason longer. Self-critique never beats majority voting and repairs about as many answers as it breaks. Debate helps only where the model cannot reason on its own: it adds +8.7 points for Gemma with thinking disabled, but +1.7 and -5.3 for models that reason, and nothing for frontier models at high reasoning effort, where all strategies tie. Letting a frontier model reason longer in one call reaches the same plateau as orchestration. The dominant factor is the question itself: across systems, question identity explains 57% of the variance in correctness, while the choice of model, strategy and reasoning effort explains 0.6%. On these tasks, additional inference, however organised, does not substitute for missing knowledge.

## 1 Introduction

LLM agents now plan experiments, analyse data and generate hypotheses in chemistry and biology [1, 2, 5, 6, 8, 15]. Their gains are typically reported against a single-call baseline, although the agent makes many calls. Because repeated sampling with a vote already improves accuracy [3, 12, 18], such comparisons conflate *what the interaction adds* with *how much inference was spent* [9].

We run a small, controlled test of this confound on biology research questions with deterministic answers. Every orchestrated strategy gets exactly the budget of the voting baseline (three calls, same model, same prompts, same decoding). We then ask a second question the agent literature rarely asks: does letting a strong model reason longer in a *single* call reach the same point?

**Findings.** (1) At matched budget, sequential critique never improves on majority voting, and multi-agent debate improves on it only for a model whose reasoning is switched off. (2) Critique is a coin flip at the level of individual answers: it rescues and corrupts at similar rates, and it is most harmful on exact sequence manipulation. (3) Debate behaves as a substitute for reasoning: for frontier models, raising reasoning effort lifts a single call to the level that debate reaches, and stacking orchestration on top adds nothing. (4) Most strikingly, whether a system is right depends overwhelmingly on *which question* it faces and barely on *how* it spends inference: every system solves the same questions and fails the same questions (Fig. 2). The remaining errors look like missing knowledge, not missing deliberation.

## 2 Setup

**Tasks.** We use four LAB-Bench categories [10] (ProtocolQA, SeqQA, DbQA, LitQA2; multiple choice with shuffled options and no abstain option) and the deterministic subset of LABBench2 SeqQA2 [11] (open-ended sequence computations with unique answers, scored by the official validators, with input files injected into the prompt). We exclude figure/table tasks (text-only models) and prompts over 6,000 characters. With a fixed seed we draw 30 questions per category (150 total), plus 10 pilot questions that are excluded from all results. No LLM judge is used anywhere.

**Table 1:** Accuracy (%) at a matched budget of three calls; DIRECT uses one call.  $\Delta$  is DEBATE – VOTE in points with a paired bootstrap 95% CI.

| Model        | $n$ | DIRECT | VOTE | CRITIQUE | DEBATE | $\Delta$ [95% CI]  |
|--------------|-----|--------|------|----------|--------|--------------------|
| GPT-OSS-120B | 118 | 46.6   | 47.5 | 44.1     | 49.2   | +1.7 [-5.9, +9.3]  |
| Gemma 4 31B  | 150 | 48.7   | 48.7 | 47.3     | 57.3   | +8.7 [+2.0, +15.3] |
| Flash-Lite   | 75  | 61.3   | 62.7 | 56.0     | 57.3   | -5.3 [-13.3, +1.3] |

**Strategies (Fig. 3 in App. A).** DIRECT makes one call. VOTE makes three independent calls and takes the majority answer, breaking ties by self-reported confidence. CRITIQUE solves, critiques the solution, then revises. DEBATE has a molecular biologist and a skeptical reviewer solve independently, and a third call adjudicates between them. VOTE, CRITIQUE and DEBATE each use exactly three calls. Only visible outputs are passed between calls.

**Models.** The main study uses GPT-OSS-120B [13] (`reasoning_effort=low`), Gemma 4 31B [16] (thinking off, the only non-default setting the API accepts), and Gemini 3.5 Flash-Lite (thinking low), all through free API tiers. Decoding uses temperature 0.6, fixed per-call seeds and a 2,048-token cap. GPT-OSS and Gemma answer all 150 questions and Flash-Lite a fixed 75-question subset. For the reasoning-effort analysis we additionally run GPT-6 Luna, GPT-6 Sol and GPT-5.6 Terra (all strategies, high effort; Luna also at low effort) and Flash-Lite (DIRECT at high thinking) on the first 30 questions in the fixed run order (6 per category).

**Pre-registration.** Task IDs, prompts, decoding, hypotheses and the analysis script were committed before the main run. The primary contrast is DEBATE – VOTE, tested with a paired bootstrap over questions (10,000 resamples) and an exact McNemar test. Post-freeze deviations are logged in the repository. They are rate-limit workarounds, a second inference provider for 21 GPT-OSS questions (each question served entirely by one provider), and the exploratory reasoning-effort arm.


### 3 Results

**Orchestration beats voting only when the model cannot think.** DEBATE improves on VOTE for Gemma, whose thinking is disabled (+8.7 points, 95% CI [+2.0, +15.3], McNemar  $p = 0.01$ ), and the gain survives excluding every question with a truncated call. For the two models that reason, the interval includes zero ( $p = 0.83$  and  $0.29$ ). CRITIQUE is never better than VOTE (−3.4, −1.3, −6.7 points). Voting itself barely improves on a single call (+0.8, +0.0, +1.3), because the three samples usually agree: only 38%, 42% and 27% of questions produce any disagreement.


**Critique repairs about as often as it breaks.** Relative to the vote, CRITIQUE rescues 27%, 13% and 25% of wrong answers but corrupts 38%, 16% and 26% of correct ones. The damage concentrates on exact sequence manipulation (SeqQA: −25, +10, −20 points vs. VOTE), where a fluent but wrong critique talks the model out of a correct computation. Protocol troubleshooting, a judgement task, is the one category where critique is neutral or positive (−4, +3, +7).

**Thinking longer reaches the same plateau.** The same pattern holds within one model. At low effort (no reasoning tokens), GPT-6 Luna gains from debate (60% vs. 53% for VOTE), just like Gemma. A single high-effort call (60%) matches low-effort DEBATE (60%) and exceeds a low-effort single call (47%). Adding orchestration on top of high effort lifts Luna only to 63–63% while generating  $\sim 3\times$  more tokens (Fig. 1). For Flash-Lite, high thinking does not help at all (67% vs. 70%) despite  $\sim 5\times$  the tokens. For GPT-6 Sol at high effort, all four strategies score exactly 70% while the three-call strategies generate  $\sim 3\times$  the tokens. Orchestration and extended reasoning are two routes to the same ceiling, not complementary gains.

**The question, not the method, determines the outcome.** A two-way decomposition of the 0/1 correctness matrix (systems  $\times$  questions) attributes 57% of the variance to the question and 0.6% to the system for the main models (Fig. 2). Across all 25 systems on the shared 30 questions, including frontier models at high reasoning, the split is 48% vs. 3.3%. The questions no system solves are dominated by literature and database recall (LitQA2, DbQA), which no amount of deliberation can supply without retrieval.



**Figure 1:** Accuracy vs. generated tokens per question on the same 30 questions. Circles are low reasoning and stars are high reasoning; colour marks the strategy. Across a  $\sim 50\times$  range of generated tokens, all ways of spending inference land on a narrow per-model plateau.



**Figure 2:** Correctness of every model  $\times$  strategy system (rows) on every question answered by all of them (columns, sorted by the fraction of systems correct). The staircase shows that systems solve and fail the *same* questions and differ mainly in a shared boundary. Question identity explains 57% [49, 64] of the variance in correctness; the system explains 0.6% [0.3, 2.2].

## 4 Discussion

On objectively scored biology research tasks, the additional calls in agentic pipelines buy little beyond what independent sampling buys, and interaction can hurt when a confident critic overrides a correct computation. For strong models, the same plateau is reached more cheaply by letting one call reason longer. The practical implication for agent builders is to measure against a compute-matched vote [9] and to invest the budget in *knowledge access* (retrieval, databases, code execution) rather than in more elaborate conversations between copies of the same model. Our results agree with compute-matched analyses of debate on general reasoning benchmarks [4, 14, 17, 19] and with evidence that intrinsic self-correction is unreliable [7]. We extend them to biology, to sequential critique, and to the trade-off between orchestration and reasoning effort.

**Limitations.** Models had no tools or retrieval, so DbQA and LitQA2 measure recall; orchestration with tools may behave differently. The budget is fixed at three calls. Frontier models were run through the Codex CLI, which adds its own system prompt and does not expose temperature or seeds; they cover 30 questions, so their comparisons are exploratory. Gemma could only run with thinking off. One prompt set per strategy was used, and 150 questions give limited power for effects under  $\sim 10$  points.

## References

- [1] Daniil A. Boiko, Robert MacKnight, Ben Kline, and Gabe Gomes. Autonomous chemical research with large language models. *Nature*, 624:570–578, 2023. URL <https://doi.org/10.1038/s41586-023-06792-0>.
- [2] Andres M. Bran, Sam Cox, Andrew D. White, and Philippe Schwaller. Chemcrow: Augmenting large-language models with chemistry tools. *arXiv preprint arXiv:2304.05376*, 2023. URL <https://arxiv.org/abs/2304.05376>.
- [3] Bradley Brown, Jordan Juravsky, Ryan Ehrlich, Ronald Clark, Quoc V. Le, Christopher Ré, and Azalia Mirhoseini. Large language monkeys: Scaling inference compute with repeated sampling. *arXiv preprint arXiv:2407.21787*, 2024. URL <https://arxiv.org/abs/2407.21787>.
- [4] H. K. Choi, X. Zhu, and S. Li. Debate or vote: Which yields better decisions in multi-agent large language models? *arXiv preprint arXiv:2508.17536*, 2025. URL <https://arxiv.org/abs/2508.17536>. NeurIPS 2025 Spotlight.
- [5] Ali E. Ghareeb, Benjamin Chang, Ludovico Mitchener, Angela Yiu, Caralyn J. Szostkiewicz, Dmytro Shved, Gavin J. Gyimesi, Jon M. Laurent, et al. A multi-agent system for automating scientific discovery. *Nature*, 2026. URL <https://doi.org/10.1038/s41586-026-10652-y>.
- [6] Juraj Gottweis, Wei-Hung Weng, Alexander Daryin, Tao Tu, Petar Sirkovic, Artiom Myaskovsky, Grzegorz Glowaty, Felix Weissenberger, et al. Accelerating scientific discovery with co-scientist. *Nature*, 2026. URL <https://doi.org/10.1038/s41586-026-10644-y>.
- [7] Jie Huang, Xinyun Chen, Swaroop Mishra, Huaixiu Steven Zheng, Adams Wei Yu, Xinying Song, and Denny Zhou. Large language models cannot self-correct reasoning yet. *arXiv preprint arXiv:2310.01798*, 2023. URL <https://arxiv.org/abs/2310.01798>. ICLR 2024.
- [8] Kexin Huang, Serena Zhang, Hanchen Wang, Yuanhao Qu, Yingzhou Lu, Yusuf Roohani, Ryan Li, Lin Qiu, et al. Biomni: A general-purpose biomedical ai agent. *bioRxiv preprint*, 2025. URL <https://www.biorxiv.org/content/10.1101/2025.05.30.656746v1>.
- [9] Sayash Kapoor, Benedikt Stroebl, Zachary S. Siegel, Nitya Nadgir, and Arvind Narayanan. Ai agents that matter. *arXiv preprint arXiv:2407.01502*, 2024. URL <https://arxiv.org/abs/2407.01502>.
- [10] Jon M. Laurent, Joseph D. Janizek, Michael Ruzo, Michaela M. Hinks, Michael J. Hammerling, Siddharth Narayanan, Manvitha Ponnampati, Andrew D. White, and Samuel G. Rodriques. Lab-bench: Measuring capabilities of language models for biology research. *arXiv preprint arXiv:2407.10362*, 2024. URL <https://arxiv.org/abs/2407.10362>.
- [11] Jon M Laurent, Albert Bou, Michael Pieler, Conor Igoe, Alex Andonian, Siddharth Narayanan, James Braza, Alexandros Sanchez Vassopoulos, et al. Labbench2: An improved benchmark for ai systems performing biology research. *arXiv preprint arXiv:2604.09554*, 2026. URL <https://arxiv.org/abs/2604.09554>.
- [12] Lingyao Li and Ze Zhang. More agents is all you need. *Transactions on Machine Learning Research*, 2024. URL <https://arxiv.org/abs/2402.05120>.
- [13] OpenAI. gpt-oss-120b & gpt-oss-20b model card. *arXiv preprint arXiv:2508.10925*, 2025. URL <https://arxiv.org/abs/2508.10925>.
- [14] Andries Smit, Nathan Grinsztajn, Paul Duckworth, Thomas D. Barrett, and Arnu Pretorius. Should we be going mad? a look at multi-agent debate strategies for llms. In *International Conference on Machine Learning*, 2024. URL <https://arxiv.org/abs/2311.17371>.
- [15] Kyle Swanson, Wesley Wu, Nash L. Bulaong, John E. Pak, and James Zou. The virtual lab: Ai agents design new sars-cov-2 nanobodies with experimental validation. *bioRxiv preprint*, 2024. URL <https://www.biorxiv.org/content/10.1101/2024.11.11.623004v1>. Published in Nature 2025.

- [16] Gemma Team, Sherif El Abd, Vaibhav Aggarwal, Robin Algayres, Alek Andreev, Olivier Bachem, Ian Ballantyne, Cormac Brick, et al. Gemma 4 technical report. *arXiv preprint arXiv:2607.02770*, 2026. URL <https://arxiv.org/abs/2607.02770>.
- [17] Qineng Wang, Zihao Wang, Ying Su, Hanghang Tong, and Yangqiu Song. Rethinking the bounds of llm reasoning: Are multi-agent discussions the key? *arXiv preprint arXiv:2402.18272*, 2024. URL <https://arxiv.org/abs/2402.18272>.
- [18] Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-consistency improves chain of thought reasoning in language models. *arXiv preprint arXiv:2203.11171*, 2022. URL <https://arxiv.org/abs/2203.11171>.
- [19] Haolun Wu, Zhenkun Li, and Lingyao Li. Can llm agents really debate? a controlled study of multi-agent debate in logical reasoning. *arXiv preprint arXiv:2511.07784*, 2025. URL <https://arxiv.org/abs/2511.07784>.

## A Prompts

All prompts are single user messages. {QUESTION} is the question (with protocol, options or injected files) and {ANSWER INSTRUCTION} is the format line for multiple choice or the official SeqQA2 suffix.

### Direct / Independent×3 / Critique call 1 (solve)

Solve the following biology research question.

{QUESTION}

{ANSWER INSTRUCTION}

Respond in exactly this format and keep it short:

<answer>YOUR ANSWER</answer>

CONFIDENCE: <probability between 0 and 1 that your answer is correct>

JUSTIFICATION: <at most two sentences>

### Critique call 2 (critique)

Below are a biology research question and a proposed solution. Act as a critical reviewer: identify any scientific, computational or reasoning error in the solution. If you think the answer is wrong, say which answer you believe is correct and why. If you find no error, say so.

{QUESTION}

Proposed solution:

{CALL 1 OUTPUT}

Respond with at most four sentences, starting with "CRITIQUE:".

### Critique call 3 (revise)

Below are a biology research question, a first attempt at solving it, and a reviewer's critique of that attempt. Either may contain mistakes. Considering both, produce the final answer.

{QUESTION}

First attempt:

{CALL 1 OUTPUT}

Reviewer's critique:

{CALL 2 OUTPUT}

{ANSWER INSTRUCTION}

Respond in exactly this format and keep it short:

<answer>YOUR ANSWER</answer>  
CONFIDENCE: <probability between 0 and 1 that your answer is correct>  
JUSTIFICATION: <at most two sentences>

### Debate call 1 (researcher)

You are a molecular biology researcher. Solve the following biology research question independently.

{QUESTION}

{ANSWER INSTRUCTION}

Respond in exactly this format and keep it short:

<answer>YOUR ANSWER</answer>  
CONFIDENCE: <probability between 0 and 1 that your answer is correct>  
JUSTIFICATION: <at most two sentences>

### Debate call 2 (skeptic)

You are a skeptical scientific reviewer. Solve the following biology research question independently, looking especially for subtle experimental, biological or computational errors and traps.

{QUESTION}

{ANSWER INSTRUCTION}

Respond in exactly this format and keep it short:

<answer>YOUR ANSWER</answer>  
CONFIDENCE: <probability between 0 and 1 that your answer is correct>  
JUSTIFICATION: <at most two sentences>

### Debate call 3 (adjudicator)

Two scientists independently analysed the biology research question below. Adjudicate between their analyses (either or both may be wrong) and submit the final answer.

{QUESTION}

Analysis 1 (molecular biology researcher):  
{CALL 1 OUTPUT}

Analysis 2 (skeptical scientific reviewer):  
{CALL 2 OUTPUT}

{ANSWER INSTRUCTION}

Respond in exactly this format and keep it short:

<answer>YOUR ANSWER</answer>  
CONFIDENCE: <probability between 0 and 1 that your answer is correct>  
JUSTIFICATION: <at most two sentences>

## B Additional results

| Category   | Gemma |      |      |      |          | Lite |      |      |      |          | OSS  |      |      |      |          |
|------------|-------|------|------|------|----------|------|------|------|------|----------|------|------|------|------|----------|
|            | D     | I    | C    | M    | Disagree | D    | I    | C    | M    | Disagree | D    | I    | C    | M    | Disagree |
| DbQA       | 56.7  | 53.3 | 40.0 | 56.7 | 33.3     | 40.0 | 40.0 | 40.0 | 33.3 | 26.7     | 39.1 | 43.5 | 47.8 | 34.8 | 34.8     |
| LitQA2     | 33.3  | 30.0 | 33.3 | 36.7 | 46.7     | 60.0 | 53.3 | 53.3 | 46.7 | 26.7     | 21.7 | 26.1 | 43.5 | 39.1 | 47.8     |
| ProtocolQA | 56.7  | 56.7 | 60.0 | 70.0 | 20.0     | 66.7 | 80.0 | 86.7 | 73.3 | 26.7     | 62.5 | 50.0 | 45.8 | 50.0 | 33.3     |
| SeqQA      | 56.7  | 56.7 | 66.7 | 73.3 | 46.7     | 73.3 | 73.3 | 53.3 | 66.7 | 13.3     | 62.5 | 66.7 | 41.7 | 62.5 | 33.3     |
| SeqQA2     | 40.0  | 46.7 | 36.7 | 50.0 | 63.3     | 66.7 | 66.7 | 46.7 | 66.7 | 40.0     | 45.8 | 50.0 | 41.7 | 58.3 | 41.7     |

**Table 2:** Per-category accuracy (D=Direct, I=Independent×3, C=Critique, M=Multi). Numbers are percentages.