
Tools, Not Talk: Agent Orchestration and Extended Reasoning Hit a Knowledge Ceiling that Tool Use Breaks on Biology Research Tasks

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Agentic systems for biology usually report gains over a single model call while
2 spending several calls, which conflates what the orchestration adds with how much
3 inference is spent. On 150 pre-registered, deterministically scored biology re-
4 search questions (LAB-Bench, LABBench2), we compare majority voting, self-
5 critique and a two-scientist debate at a matched budget of three calls across three
6 models, and we contrast them with longer reasoning and with tool use. Critique
7 never beats voting, and debate helps only a model whose reasoning is switched
8 off (+8.7 points). Raising reasoning effort rarely helps (GPT-OSS spends $25\times$
9 more tokens for no gain). Instead, correctness is set by the question: question
10 identity explains 56% of its variance and the choice of model, strategy and effort
11 0.6%. Giving the model tools breaks this ceiling. A single tool-using call with
12 web search and code execution raises accuracy from 47% to 70%, and solves 47%
13 of the questions that no orchestrated or reasoning-heavy system solved (vs. 6%
14 without tools). For life-science agents, the budget belongs in tools, not in talk.

15 1 Introduction

16 LLM agents now plan experiments, analyse data and generate hypotheses in chemistry and biology
17 [1, 2, 5, 6, 8, 15]. Their gains are typically reported against a single-call baseline, although the agent
18 makes many calls. Because repeated sampling with a vote already improves accuracy [3, 12, 18],
19 such comparisons conflate *what the interaction adds* with *how much inference was spent* [9].

20 For anyone building life-science agents, this is a design question as much as an evaluation one:
21 should the budget go into richer orchestration (critics, debating scientists) or somewhere else? We
22 run a small, controlled test of this confound on biology research questions with deterministic an-
23 swers. Every orchestrated strategy gets exactly the budget of the voting baseline (three calls, same
24 model, same prompts, same decoding). We then ask a second question the agent literature rarely
25 asks: does letting a strong model reason longer in a *single* call reach the same point?

26 **Findings.** (1) At matched budget, sequential critique never improves on majority voting, and
27 multi-agent debate improves on it only for a model whose reasoning is switched off. (2) Critique
28 never helps: it breaks at least as many correct answers as it repairs, and for models that reason it is
29 most harmful on exact sequence manipulation. (3) Thinking longer is not a way out: higher reason-
30 ing effort helps only a model that otherwise does not reason at all, and then only up to the level its
31 low-effort debate already reached; stacking orchestration on top adds nothing reliable. (4) Whether
32 a system is right depends overwhelmingly on *which question* it faces and barely on *how* it spends
33 inference: every system solves and fails the same questions (Fig. 2). (5) These errors are missing

Table 1: Accuracy (%) at a matched budget of three calls; DIRECT uses one call. Δ is DEBATE – VOTE in points with a paired bootstrap 95% CI.

Model	n	DIRECT	VOTE	CRITIQUE	DEBATE	Δ [95% CI]
GPT-OSS-120B	150	46.7	49.3	42.7	47.3	-2.0 [-8.7, +4.7]
Gemma 4 31B	150	48.7	48.7	47.3	57.3	+8.7 [+2.0, +15.3]
Flash-Lite	75	61.3	62.7	56.0	57.3	-5.3 [-13.3, +1.3]

34 knowledge, not missing deliberation. One tool-using call with web search and code execution solves
 35 nearly half of the questions that every other configuration failed.

36 2 Setup

37 **Tasks.** We use four LAB-Bench categories [10] (ProtocolQA, SeqQA, DbQA, LitQA2; multiple
 38 choice with shuffled options and no abstain option) and the deterministic subset of LABBench2
 39 SeqQA2 [11] (open-ended sequence computations with unique answers, scored by the official val-
 40 idators, with input files injected into the prompt). We exclude figure/table tasks (text-only models)
 41 and prompts over 6,000 characters. With a fixed seed we draw 30 questions per category (150 total),
 42 plus 10 pilot questions that are excluded from all results. No LLM judge is used anywhere.

43 **Strategies (Fig. 3 in App. A).** DIRECT makes one call. VOTE makes three independent calls and
 44 takes the majority answer, breaking ties by self-reported confidence. CRITIQUE solves, critiques the
 45 solution, then revises. DEBATE has a molecular biologist and a skeptical reviewer solve indepen-
 46 dently, and a third call adjudicates between them. VOTE, CRITIQUE and DEBATE each use exactly
 47 three calls. Only visible outputs are passed between calls.

48 **Models.** The main study uses GPT-OSS-120B [13] (reasoning_effort=low), Gemma 4 31B
 49 [16] (thinking off, the only non-default setting the API accepts), and Gemini 3.5 Flash-Lite (thinking
 50 low), all through free API tiers. Decoding uses temperature 0.6, fixed per-call seeds and a 2,048-
 51 token cap. GPT-OSS and Gemma answer all 150 questions and Flash-Lite a fixed 75-question
 52 subset. For the reasoning-effort analysis we additionally run GPT-6 Luna, GPT-6 Sol and GPT-5.6
 53 Terra (all strategies, high effort; Luna also at low effort), GPT-OSS at high effort (all strategies,
 54 served by Ollama; a low-effort DIRECT run on the same host matches the main run’s accuracy)
 55 and Flash-Lite (DIRECT at high thinking), all on the first 30 questions in the fixed run order (6 per
 56 category).

57 **Pre-registration.** Task IDs, prompts, decoding, hypotheses and the analysis script were commit-
 58 ted before the main run. The primary contrast is DEBATE – VOTE, tested with a paired bootstrap
 59 over questions (10,000 resamples) and an exact McNemar test. Post-freeze deviations are listed
 60 in App. B: rate-limit workarounds, a second inference provider for 22 GPT-OSS questions (each
 61 question served entirely by one provider), and the exploratory reasoning-effort arm.

62 3 Results

63 **Orchestration beats voting only when the model cannot think.** DEBATE improves on VOTE for
 64 Gemma, whose thinking is disabled (+8.7 points, 95% CI [+2.0, +15.3], McNemar $p = 0.015$),
 65 and the gain survives excluding every question with a truncated call. For the two models that reason,
 66 the interval includes zero ($p = 0.7$ and 0.29 ; for GPT-OSS the estimate is unchanged at -1.6 when
 67 the 22 questions served by the second provider are excluded). Pooled over all 375 model–question
 68 pairs, DEBATE – VOTE is $+1.6$ [$-2.4, +5.6$] points and CRITIQUE – VOTE is -4.5 [$-9.3, +0.3$].
 69 CRITIQUE is never better than VOTE ($-6.7, -1.3, -6.7$ points). Voting itself barely improves on
 70 a single call ($+2.7, +0.0, +1.3$ points), even though the three samples disagree on 38%, 42% and
 71 27% of questions.

72 **Critique breaks at least as often as it repairs.** Relative to the vote, CRITIQUE rescues 24%, 13%
 73 and 25% of wrong answers but corrupts 38%, 16% and 26% of correct ones. For the two models




Figure 1: Accuracy vs. generated tokens per question on the same 30 questions. Circles are low reasoning and stars are high reasoning; colour marks the strategy. Across nearly two orders of magnitude in generated tokens, all ways of spending inference land on a narrow per-model plateau; the black diamond (GPT-6 Luna, one call with web search and code execution) is the only configuration that rises above it, at a fraction of the tokens.




Figure 2: Correctness of every model $\times$ strategy system (rows) on every question answered by all of them (columns, sorted by the fraction of systems correct). The staircase shows that systems solve and fail the *same* questions and differ mainly in a shared boundary. Question identity explains 56% [49, 62] of the variance in correctness; the system explains 0.6% [0.3, 1.9].

74 that reason, the largest losses are on exact sequence manipulation (SeqQA: -23 and -20 points vs.
75 VOTE for GPT-OSS and Flash-Lite; +10 for Gemma), where a fluent but wrong critique talks the
76 model out of a correct computation.

77 **Thinking longer mostly does not help either.** On the same 30 questions (Fig. 1), raising GPT-
78 OSS to high reasoning effort multiplies generated tokens by $25\times$ (284 to 7,191 per question) while
79 accuracy stays flat (50% to 47%). For Flash-Lite the result is the same (70% to 67%, $5\times$ tokens).
80 The one model that gains is GPT-6 Luna, which at low effort uses no reasoning tokens (47% to
81 60%). Its high-effort single call lands exactly where its low-effort DEBATE already was (60%),
82 consistent with debate acting as a substitute for reasoning. Orchestration on top of high effort adds
83 nothing reliable. GPT-6 Sol scores 70% with all four strategies, and GPT-5.6 Terra ranges from 53%
84 (CRITIQUE) to 67% (DEBATE) with no paired difference significant. Within each model, accuracy
85 moves by a few questions while generated tokens vary by up to two orders of magnitude.

86 **The question, not the method, determines the outcome.** A two-way decomposition of the 0/1
87 correctness matrix (systems $\times$ questions) attributes 56% of the variance to the question and 0.6%
88 to the system for the main models (Fig. 2). Across all 33 systems on the shared 30 questions,
89 including frontier models at high reasoning, the split is 47% vs. 3.9%. Of the 32 questions that no
90 main system solves, 18 require literature or database recall (LitQA2, DbQA) and 9 are open-ended
91 sequence computations (SeqQA2), failures that point to missing retrieval and code execution rather
92 than missing deliberation.

93 **Tools break the ceiling.** To test whether the ceiling reflects missing knowledge, we give GPT-6
94 Luna (low effort, one call) a shell with Python and live web search, announced by a one-sentence
95 harness note, and leave everything else unchanged. On the shared 30 questions accuracy rises from
96 47% to 70% (+23 points, 95% CI [+7, +43], McNemar $p = 0.039$), more than any orchestration or
97 reasoning setting, while generating only ~ 404 tokens per question. On the 32 questions that none of

the eight main systems solved, it reaches 47% vs. 6% for the same model without tools (+41 points [+25, +56], $p = < 0.001$). The gain comes from retrieval: LitQA2 rises from 18% to 91% and DbQA from 0% to 43% on these questions, whereas code execution barely moves the open-ended SeqQA2 computations (0% to 11%). All 77 searches and 19 commands were logged, and none reached the benchmark or its datasets.

4 Implications for building life-science agents

Benchmark against a compute-matched vote. A multi-call agent should beat independent sampling with the same number of calls, not a single call [9]. Most orchestration gains vanish under that baseline. **Spend the harness budget on tools, not talk.** Critics, debaters and longer reasoning re-deliberate over the same missing knowledge. Retrieval over literature and databases supplies it, at a fraction of the tokens. **Sequence computation needs better tools than a bare shell.** Generic code execution rarely fixed the open-ended sequence tasks, which suggests that domain tool ecosystems (sequence libraries, primer and cloning tools) and specialist models are where the next gains lie. Our results agree with compute-matched analyses of debate on general reasoning benchmarks [4, 14, 17, 19] and with evidence that intrinsic self-correction is unreliable [7]. We extend them to biology, to sequential critique, and to the trade-off between orchestration and reasoning effort.

Limitations. Models had no tools or retrieval, so DbQA and LitQA2 measure recall; orchestration with tools may behave differently. The budget is fixed at three calls. Frontier models were run through the Codex CLI, which adds its own system prompt and does not expose temperature or seeds; they cover 30 questions, so their comparisons are exploratory. Gemma could only run with thinking off. One prompt set per strategy was used, and 150 questions give limited power for effects under ~ 10 points. The tools arm uses one model, a harness note that the no-tools condition lacks, and 30+32 questions; web retrieval can in principle surface benchmark-derived pages, which our logs did not show.

5 Conclusion

On objectively scored biology research tasks, how a model spends extra inference (voting, self-critique, multi-agent debate or longer reasoning) barely changes which questions it gets right: question identity explains 56% of the variance in correctness and the choice of system 0.6%. The ceiling is knowledge, not deliberation, and a single tool-using call breaks it where no amount of orchestration did. Builders of life-science agents should spend their budget on tools, not talk.

References

- [1] Daniil A. Boiko, Robert MacKnight, Ben Kline, and Gabe Gomes. Autonomous chemical research with large language models. *Nature*, 624:570–578, 2023. URL <https://doi.org/10.1038/s41586-023-06792-0>.
- [2] Andres M. Bran, Sam Cox, Andrew D. White, and Philippe Schwaller. Chemcrow: Augmenting large-language models with chemistry tools. *arXiv preprint arXiv:2304.05376*, 2023. URL <https://arxiv.org/abs/2304.05376>.
- [3] Bradley Brown, Jordan Juravsky, Ryan Ehrlich, Ronald Clark, Quoc V. Le, Christopher Ré, and Azalia Mirhoseini. Large language monkeys: Scaling inference compute with repeated sampling. *arXiv preprint arXiv:2407.21787*, 2024. URL <https://arxiv.org/abs/2407.21787>.
- [4] H. K. Choi, X. Zhu, and S. Li. Debate or vote: Which yields better decisions in multi-agent large language models? *arXiv preprint arXiv:2508.17536*, 2025. URL <https://arxiv.org/abs/2508.17536>. NeurIPS 2025 Spotlight.
- [5] Ali E. Ghareeb, Benjamin Chang, Ludovico Mitchener, Angela Yiu, Caralyn J. Szostkiewicz, Dmytro Shved, Gavin J. Gyimesi, Jon M. Laurent, et al. A multi-agent system for automating scientific discovery. *Nature*, 2026. URL <https://doi.org/10.1038/s41586-026-10652-y>.

- [6] Juraj Gottweis, Wei-Hung Weng, Alexander Daryin, Tao Tu, Petar Sirkovic, Artiom Myaskovsky, Grzegorz Glowaty, Felix Weissenberger, et al. Accelerating scientific discovery with co-scientist. *Nature*, 2026. URL <https://doi.org/10.1038/s41586-026-10644-y>.
- [7] Jie Huang, Xinyun Chen, Swaroop Mishra, Huaixiu Steven Zheng, Adams Wei Yu, Xinying Song, and Denny Zhou. Large language models cannot self-correct reasoning yet. *arXiv preprint arXiv:2310.01798*, 2023. URL <https://arxiv.org/abs/2310.01798>. ICLR 2024.
- [8] Kexin Huang, Serena Zhang, Hanchen Wang, Yuanhao Qu, Yingzhou Lu, Yusuf Roohani, Ryan Li, Lin Qiu, et al. Biomni: A general-purpose biomedical ai agent. *bioRxiv preprint*, 2025. URL <https://www.biorxiv.org/content/10.1101/2025.05.30.656746v1>.
- [9] Sayash Kapoor, Benedikt Stroebel, Zachary S. Siegel, Nitya Nadgir, and Arvind Narayanan. Ai agents that matter. *arXiv preprint arXiv:2407.01502*, 2024. URL <https://arxiv.org/abs/2407.01502>.
- [10] Jon M. Laurent, Joseph D. Janizek, Michael Ruzo, Michaela M. Hinks, Michael J. Hammerling, Siddharth Narayanan, Manvitha Ponnampati, Andrew D. White, and Samuel G. Rodrigues. Lab-bench: Measuring capabilities of language models for biology research. *arXiv preprint arXiv:2407.10362*, 2024. URL <https://arxiv.org/abs/2407.10362>.
- [11] Jon M Laurent, Albert Bou, Michael Pieler, Conor Igoe, Alex Andonian, Siddharth Narayanan, James Braza, Alexandros Sanchez Vassopoulos, et al. Labbench2: An improved benchmark for ai systems performing biology research. *arXiv preprint arXiv:2604.09554*, 2026. URL <https://arxiv.org/abs/2604.09554>.
- [12] Lingyao Li and Ze Zhang. More agents is all you need. *Transactions on Machine Learning Research*, 2024. URL <https://arxiv.org/abs/2402.05120>.
- [13] OpenAI. gpt-oss-120b & gpt-oss-20b model card. *arXiv preprint arXiv:2508.10925*, 2025. URL <https://arxiv.org/abs/2508.10925>.
- [14] Andries Smit, Nathan Grinsztajn, Paul Duckworth, Thomas D. Barrett, and Arnau Pretorius. Should we be going mad? a look at multi-agent debate strategies for llms. In *International Conference on Machine Learning*, 2024. URL <https://arxiv.org/abs/2311.17371>.
- [15] Kyle Swanson, Wesley Wu, Nash L. Bulaong, John E. Pak, and James Zou. The virtual lab: Ai agents design new sars-cov-2 nanobodies with experimental validation. *bioRxiv preprint*, 2024. URL <https://www.biorxiv.org/content/10.1101/2024.11.11.623004v1>. Published in Nature 2025.
- [16] Gemma Team, Sherif El Abd, Vaibhav Aggarwal, Robin Algayres, Alek Andreev, Olivier Bachem, Ian Ballantyne, Cormac Brick, et al. Gemma 4 technical report. *arXiv preprint arXiv:2607.02770*, 2026. URL <https://arxiv.org/abs/2607.02770>.
- [17] Qineng Wang, Zihao Wang, Ying Su, Hanghang Tong, and Yangqiu Song. Rethinking the bounds of llm reasoning: Are multi-agent discussions the key? *arXiv preprint arXiv:2402.18272*, 2024. URL <https://arxiv.org/abs/2402.18272>.
- [18] Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-consistency improves chain of thought reasoning in language models. *arXiv preprint arXiv:2203.11171*, 2022. URL <https://arxiv.org/abs/2203.11171>.
- [19] Haolun Wu, Zhenkun Li, and Lingyao Li. Can llm agents really debate? a controlled study of multi-agent debate in logical reasoning. *arXiv preprint arXiv:2511.07784*, 2025. URL <https://arxiv.org/abs/2511.07784>.

A Prompts

All prompts are single user messages. {QUESTION} is the question (with protocol, options or injected files) and {ANSWER INSTRUCTION} is the format line for multiple choice or the official SeqQA2 suffix.

195 **Direct / Independent×3 / Critique call 1 (solve)**

196 Solve the following biology research question.

197

198 {QUESTION}

199

200 {ANSWER INSTRUCTION}

201

202 Respond in exactly this format and keep it short:

203 <answer>YOUR ANSWER</answer>

204 CONFIDENCE: <probability between 0 and 1 that your answer is correct>

205 JUSTIFICATION: <at most two sentences>

206 **Critique call 2 (critique)**

207 Below are a biology research question and a proposed solution. Act as a critical reviewer: identify any
208 scientific, computational or reasoning error in the solution. If you think the answer is wrong, say which
209 answer you believe is correct and why. If you find no error, say so.

210

211 {QUESTION}

212

213 Proposed solution:

214 {CALL 1 OUTPUT}

215

216 Respond with at most four sentences, starting with "CRITIQUE:".

217 **Critique call 3 (revise)**

218 Below are a biology research question, a first attempt at solving it, and a reviewer's critique of that
219 attempt. Either may contain mistakes. Considering both, produce the final answer.

220

221 {QUESTION}

222

223 First attempt:

224 {CALL 1 OUTPUT}

225

226 Reviewer's critique:

227 {CALL 2 OUTPUT}

228

229 {ANSWER INSTRUCTION}

230

231 Respond in exactly this format and keep it short:

232 <answer>YOUR ANSWER</answer>

233 CONFIDENCE: <probability between 0 and 1 that your answer is correct>

234 JUSTIFICATION: <at most two sentences>

235 **Debate call 1 (researcher)**

236 You are a molecular biology researcher. Solve the following biology research question independently.

237

238 {QUESTION}

239

240 {ANSWER INSTRUCTION}

241

242 Respond in exactly this format and keep it short:

243 <answer>YOUR ANSWER</answer>

244 CONFIDENCE: <probability between 0 and 1 that your answer is correct>

245 JUSTIFICATION: <at most two sentences>

246 **Debate call 2 (skeptic)**

247 You are a skeptical scientific reviewer. Solve the following biology research question independently, looking
248 especially for subtle experimental, biological or computational errors and traps.

249

250 {QUESTION}

251

252 {ANSWER INSTRUCTION}

253
 254 Respond in exactly this format and keep it short:
 255 <answer>YOUR ANSWER</answer>
 256 CONFIDENCE: <probability between 0 and 1 that your answer is correct>
 257 JUSTIFICATION: <at most two sentences>

258 **Debate call 3 (adjudicator)**

259 Two scientists independently analysed the biology research question below. Adjudicate between their analyses
 260 (either or both may be wrong) and submit the final answer.
 261
 262 {QUESTION}
 263
 264 Analysis 1 (molecular biology researcher):
 265 {CALL 1 OUTPUT}
 266
 267 Analysis 2 (skeptical scientific reviewer):
 268 {CALL 2 OUTPUT}
 269
 270 {ANSWER INSTRUCTION}
 271
 272 Respond in exactly this format and keep it short:
 273 <answer>YOUR ANSWER</answer>
 274 CONFIDENCE: <probability between 0 and 1 that your answer is correct>
 275 JUSTIFICATION: <at most two sentences>

276 B Pre-registration deviations

277 All task IDs, prompts, decoding settings and the primary analysis were frozen after a 10-question
 278 pilot and before any main-set call. Changes made after the freeze: (i) throttle and retry fixes for an
 279 undocumented 150 requests/hour provider limit (no completed call was re-run); (ii) 22 GPT-OSS
 280 questions served in full by a second provider of the same model (identical prompts, seeds and cap;
 281 excluding them leaves DEBATE – VOTE at -1.6), with partially served questions discarded and
 282 re-run on the primary provider; (iii) the exploratory reasoning-effort arm on the first 30 questions,
 283 including frontier models served through the Codex CLI with web search, code execution and sub-
 284 agents disabled (a first attempt with web search enabled was discarded unanalysed), and GPT-OSS
 285 at high effort on a second host.

286 C Additional results

Category	Gemma					Lite					OSS				
	D	I	C	M	Disagree	D	I	C	M	Disagree	D	I	C	M	Disagree
DbQA	56.7	53.3	40.0	56.7	33.3	40.0	40.0	40.0	33.3	26.7	46.7	46.7	43.3	36.7	36.7
LitQA2	33.3	30.0	33.3	36.7	46.7	60.0	53.3	53.3	46.7	26.7	20.0	30.0	40.0	36.7	43.3
ProtocolQA	56.7	56.7	60.0	70.0	20.0	66.7	80.0	86.7	73.3	26.7	56.7	50.0	43.3	46.7	33.3
SeqQA	56.7	56.7	66.7	73.3	46.7	73.3	73.3	53.3	66.7	13.3	63.3	66.7	43.3	60.0	33.3
SeqQA2	40.0	46.7	36.7	50.0	63.3	66.7	66.7	46.7	66.7	40.0	46.7	53.3	43.3	56.7	43.3

Table 2: Per-category accuracy (D=Direct, I=Independent×3, C=Critique, M=Multi). Numbers are percentages.